

Figure 5.6

Methods and provenance

Study boundary

22 documents · 140 chunks · 34 owner-approved questions · five systems · six categories

Two documents supply 65/140 chunks: corpus concentration limits external validity.

Comparison

Vector-free strategies: BM25, Entity-Co-occurrence Graph, Hybrid RRF, Prompt-RAG reranker.

Dense-vector baseline: FAISS windowed-max. Prompt-RAG reranks frozen BM25 top-50; it is not a generator.

Evidence contract

Primary qrels: final_pooled owner-adjudicated · bootstrap: 10,000 whole-query samples, seed 42.

H1-H4 read from preregistered seed-42 results. H5 uses real aggregate correlations from frozen evaluation.

Invalid seed-123 statistics and INVALID-seed123 qrels were blocked before read.

Interpretation

Exploratory pilot only. No composite metric, universal winner, deployed router, or causal H5 claim.

Screen 1 is the citable record; Screen 2 is a viva exhibit and is not cited as evidence.

Frozen provenance block for every dashboard figure. Source hashes and exact reproduction commands are recorded in dashboards/data/gate1_reconciliation.json and dashboards/RESULTS_LOG.md.